

Linear Regression

Minor in AI, IIT Ropar

2nd April, 2025

Introduction

Linear regression is a fundamental supervised learning technique used to predict a continuous output variable (Y) based on one or more input variables (X). It is akin to teaching a child how to recognize patterns in labeled data, much like learning to identify fruits in a supermarket. This case study explores linear regression concepts, its mathematical foundation, and practical implementation using Python.

The Curious Kid and Linear Regression

Imagine teaching a child about the relationship between age and height. You provide labeled examples: "At age 1, height is 75 cm; at age 2, height is 85 cm." The child learns that as age increases, so does height. This mirrors the essence of linear regression: finding the best-fit line that explains the relationship between input (X) and output (Y).

Core Concepts

Mathematical Foundation

The linear regression model is represented as:

$$Y = MX + B$$

where:

- M : Slope of the line (rate of change of Y with respect to X),
- B : Intercept (value of Y when $X = 0$).

Model Training

The goal is to minimize the error between actual (Y) and predicted (\hat{Y}) values by adjusting M and B .

Visualization

Scatter plots display data points. The regression line represents the relationship learned by the model.

Handling Outliers

Outliers can distort the slope and intercept, leading to inaccurate predictions. Tools like box plots help identify such anomalies.

Implementation in Python

Example 1: Predicting Salary Based on Experience

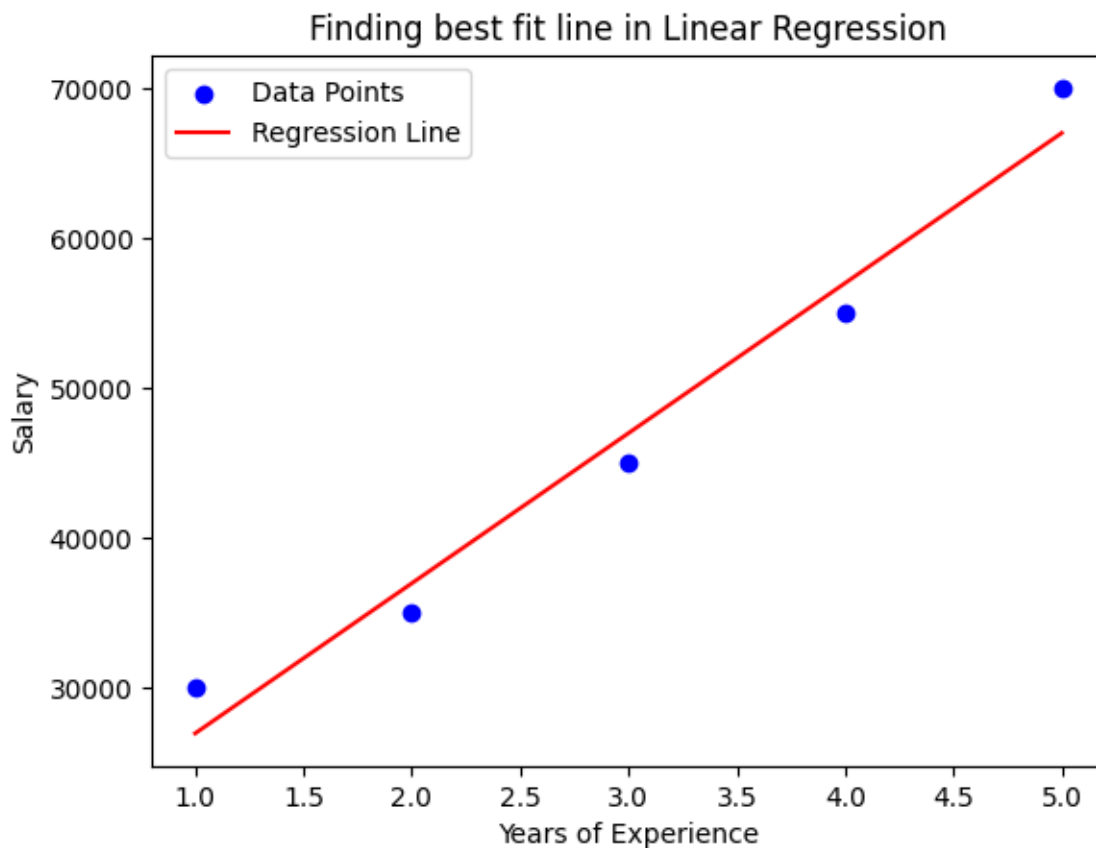
The following Python code demonstrates how to train a linear regression model using Scikit-learn:

```
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt

# Input (Years of Experience) and Output (Salary)
X = [[1], [2], [3], [4], [5]]
Y = [30000, 35000, 45000, 55000, 70000]

# Train the model
model = LinearRegression().fit(X, Y)

# Visualization
plt.scatter(X, Y, color='blue', label="Data Points")
plt.plot(X, model.predict(X), color='red', label="Regression Line")
plt.xlabel("Years of Experience")
plt.ylabel("Salary")
plt.legend()
plt.title("Salary Prediction using Linear Regression")
plt.show()
```



The plot illustrates the linear relationship between years of experience and salary.

Example 2: Predicting Outputs for New Inputs

The following Python code predicts salary for new inputs:

```
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt

# Input (Years of Experience) and Output (Salary)
X = [[1], [2], [3], [4], [5]]
Y = [30000, 35000, 45000, 55000, 70000]


# Train the model
model = LinearRegression().fit(X, Y)

# Visualization
plt.scatter(X, Y, color='blue', label="Data Points")
plt.plot(X, model.predict(X), color='red', label="Regression Line")
plt.xlabel("Years of Experience")
plt.ylabel("Salary")
plt.legend()
plt.title("Finding best fit line in Linear Regression")

# Predicting for new inputs
new_inputs = [[6], [7], [8]]
predictions = model.predict(new_inputs)

# Plotting for new inputs
plt.scatter(new_inputs, predictions, color='green', marker='x', label="Predictions")
plt.legend()
plt.show()

for inp, pred in zip(new_inputs, predictions):
    print(f"For input {inp[0]}, predicted output is {pred:.2f}")
```



This demonstrates how predictions can be made for unseen data after training, i.e. after the best fit line has been created by our linear regression model.

How Changing Slope (M) and Intercept (B) Affects the Regression Line

The slope (M) and intercept (B) fundamentally shape a regression line's behavior. Here's how they influence predictions:

1. Impact of Intercept (B)

- The intercept determines where the line crosses the Y-axis (when $X = 0$).
- Increasing B shifts the entire line **vertically upward** without changing its steepness.
- Decreasing B shifts the line **downward**, maintaining the same rate of change.

The following Python code visualizes how changes in B affect the regression line:

```
import numpy as np
import matplotlib.pyplot as plt

# Define different values of b (intercept) with a fixed m
b_values = [0, 2, 5] # Different intercepts
m = 1 # Fixed slope
X = np.linspace(-5, 5, 100) # Generate values from -5 to 5
```

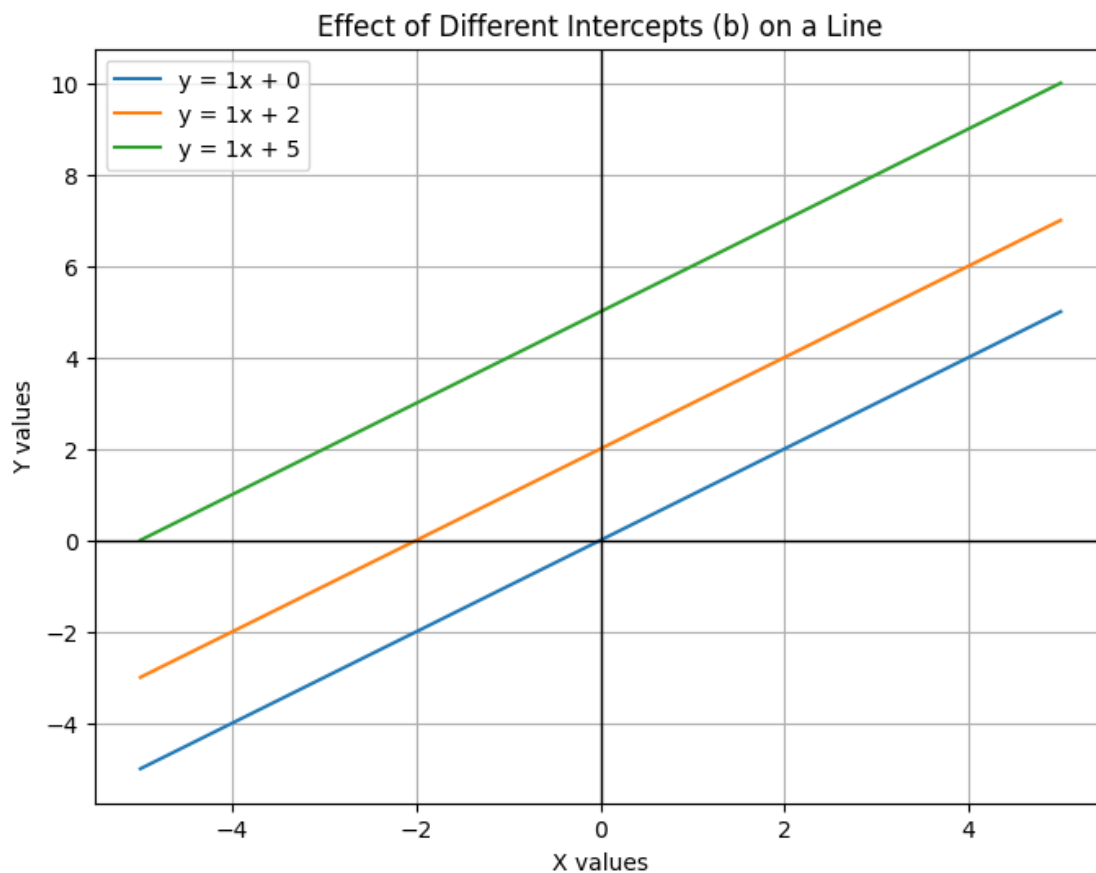
```

# Plot different intercepts
plt.figure(figsize=(8,6))
for b in b_values:
    Y = m * X + b
    plt.plot(X, Y, label=f"y = {m}x + {b}")

# Highlight origin and axes
plt.axhline(0, color='black', linewidth=1)
plt.axvline(0, color='black', linewidth=1)

# Labels and legend
plt.xlabel("X values")
plt.ylabel("Y values")
plt.title("Effect of Different Intercepts (b) on a Line")
plt.legend()
plt.grid()
plt.show()

```



The plot shows that lines shift vertically while maintaining their slope when B changes.

2. Impact of Slope (M)

- The slope defines the line's steepness and direction (positive/negative relationship).
- A larger $|M|$ creates a steeper line, indicating stronger rate of change.
- $M = 0$ produces a horizontal line (no relationship between X and Y).

The following Python code visualizes how changes in M affect the regression line:

```

import numpy as np
import matplotlib.pyplot as plt

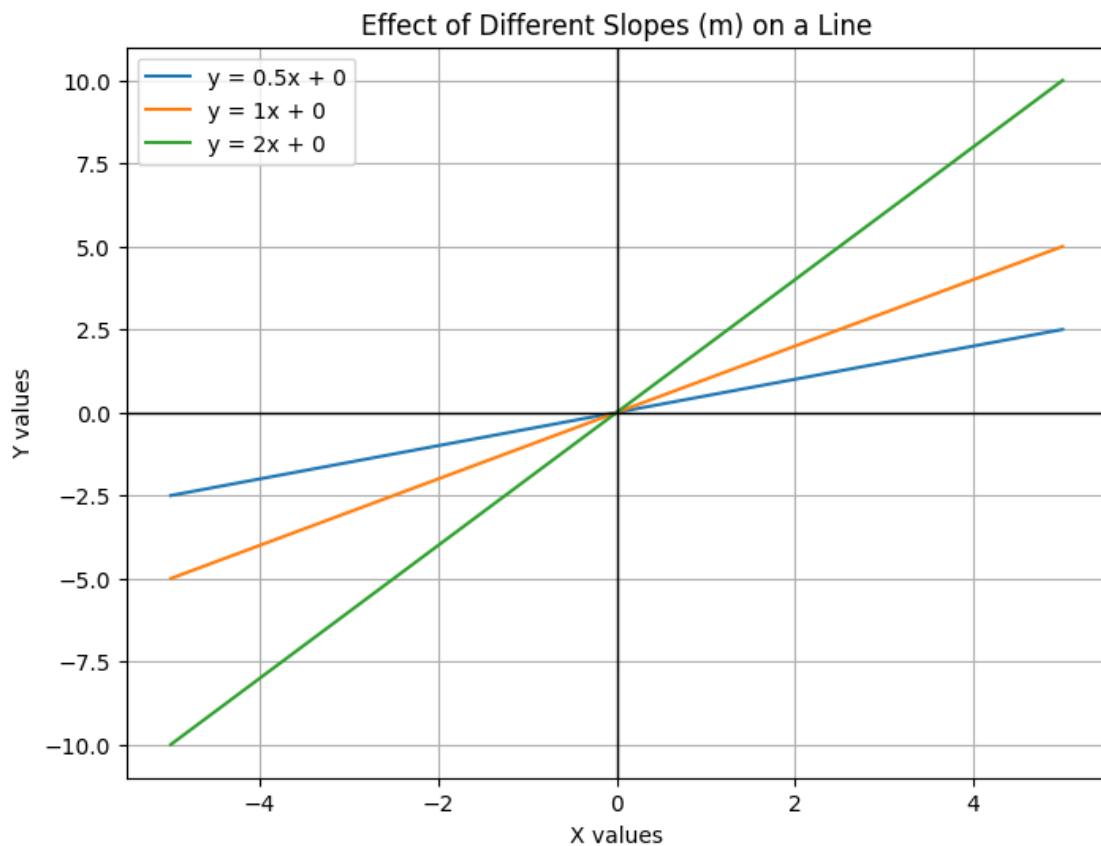
# Define different values of m (slope) with a fixed b
m_values = [0.5, 1, 2] # Different slopes
b = 0 # Fixed intercept
X = np.linspace(-5, 5, 100)

# Plot different slopes
plt.figure(figsize=(8,6))
for m in m_values:
    Y = m * X + b
    plt.plot(X, Y, label=f"y = {m}x + {b}")

# Highlight origin and axes
plt.axhline(0, color='black', linewidth=1)
plt.axvline(0, color='black', linewidth=1)

# Labels and legend
plt.xlabel("X values")
plt.ylabel("Y values")
plt.title("Effect of Different Slopes (m) on a Line")
plt.legend()
plt.grid()
plt.show()

```



The plot shows that lines rotate while maintaining their intercept when M changes.

Real-World Analogy: Supervised Learning

Linear regression aligns with supervised learning principles:

- **Labeled Data:** The model learns from examples where both inputs (X) and outputs (Y) are known.
- **Prediction:** Once trained, the model predicts outputs for new inputs.

Applications

Linear regression has widespread applications:

- Predicting house prices based on features like size or location.
- Forecasting sales trends over time.
- Estimating salaries based on years of experience.

Conclusion

Linear regression bridges mathematical theory and practical application, offering insights into relationships within data. By understanding its fundamentals and implementing it step-by-step, we can solve real-world problems effectively.