

**Minor in AI**

**Support Vector Machines Demystified**

# 1 From Gardens to Algorithms: Iris Flower

Imagine you're handed three types of iris flowers (Setosa, Versicolor, Virginica) and asked to classify them using only two measurements: **sepal length** and **sepal width**. This real-world problem mirrors the classic Iris dataset, where machines learn to distinguish species through pattern recognition.

## Why SVM?

Traditional rulers fail when data isn't linearly separable. SVMs act like **smart flexible rulers** that can bend their decision boundaries using kernel tricks. The RBF (Radial Basis Function) kernel is particularly powerful—it transforms the 2D plane into a higher-dimensional space where separation becomes possible.

## 2 The Science Behind the Magic

### 2.1 Gamma: The Neighborhood Watch Parameter

In the context of the RBF kernel,  $\gamma$  determines how “local” or “global” the influence of each training point is when shaping the decision boundary.

- **High gamma** ( $\gamma = 1.0$ ) – *Microscope mode* Each point's influence is sharply localized. Only very close neighbors contribute meaningfully to the kernel value, leading to a wiggly boundary that can overfit small fluctuations.
- **Low gamma** ( $\gamma = 0.01$ ) – *Satellite view* Each point's influence spreads broadly across the feature space. Distant points still have non-negligible effect, producing a smoother, more global decision surface.

**Real-World Analogy:** Drop a pebble in a pond and observe the ripples. Gamma controls how far those ripples travel. A high gamma is like a thick, viscous fluid where ripples die out quickly (only very close points “feel” each other). A low gamma is like water where ripples travel far, connecting distant regions.

$$\text{Influence between two points} \propto e^{-\gamma \times (\text{distance})^2}$$

Here, “distance” is the straight-line (Euclidean) separation between two points in your feature space.

### 2.2 RBF Kernel Deep Dive

The Radial Basis Function (RBF) kernel computes similarity  $K(x_i, x_j)$  between any two feature vectors  $x_i$  and  $x_j$  as:

$$K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$$

Let us unpack each component:

- $\|x_i - x_j\|^2$ : The squared Euclidean distance

$$\|x_i - x_j\|^2 = (x_{i1} - x_{j1})^2 + (x_{i2} - x_{j2})^2 + \dots$$

This measures how far apart the two points are in the original space.

- $-\gamma$ : A scaling factor in the exponent A larger  $\gamma$  makes the exponent more negative for a given distance, causing the similarity to drop off more rapidly.
- $\exp(\cdot)$ : The exponential function Transforms the negative distance into a similarity score between 0 and 1:

$$K(x_i, x_j) \rightarrow \begin{cases} 1, & \text{if } x_i = x_j \quad (\text{zero distance}) \\ 0, & \text{as } \|x_i - x_j\| \rightarrow \infty \end{cases}$$

**Interpretation:** The RBF kernel effectively creates a high-dimensional feature space where each original point  $x_i$  is represented by a “bump” centered at that point. The height of the bump at another point  $x_j$  is exactly  $K(x_i, x_j)$ . SVM then finds a linear hyperplane in this transformed space, which corresponds to a non-linear boundary back in the original space.

With an appropriate choice of  $\gamma$  (and regularization  $C$ ), the RBF kernel equips SVM with the flexibility to carve out complex decision regions—much like a sculptor shaping clay into intricate forms.

#### Spreadsheet Insight

In our lecture demo, we calculated pairwise distances between 3 points:

Points	Squared Distance	RBF Value (=0.1)
A-B	2.0	$e^{-0.1*2.0} \approx 0.8187$
A-C	5.0	$e^{-0.1*5.0} \approx 0.6065$

This creates a **similarity matrix** that guides boundary creation.

## 3 Hands-On: Classifying Flowers

Listing 1: Iris Classification Code


```

1 from sklearn import datasets
2 from sklearn.svm import SVC
3 from sklearn.preprocessing import StandardScaler
4
5 # Load data
6 iris = datasets.load_iris()
7 X = iris.data[:, :2] # Sepal features only
8 y = iris.target
9
10 # Standardization is crucial!
11 scaler = StandardScaler()
12 X_scaled = scaler.fit_transform(X)
13
14 # Build RBF-SVM model
15 model = SVC(kernel='rbf', gamma='auto')
16 model.fit(X_scaled, y)
17
18 # Visualize decision boundaries
19 plot_decision_boundary(X_scaled, y, model)

```

### Step-by-Step Explanation:

1. **Feature Selection:** Using only 2 features (sepal length/width) enables 2D visualization
2. **Standardization:** SVMs are distance-based—scaling prevents dominance by large-valued features
3. **Gamma='auto':** Lets sklearn choose  $\gamma = 1/(\text{n\_features} \times \text{X.var}())$
4. **Decision Boundary Plot:** Shows regions where model predicts each class



Complex RBF boundaries separating three iris classes

## 4 When Penguins Challenge Machines

The Palmer Penguins dataset introduces new real-world hurdles for our SVM classifier:

- **Missing values:** We must clean the data before training.
- **Overlapping classes:** Different species may share similar measurements.
- **Biological variation:** Natural variability requires robust decision boundaries.

Listing 2: Penguin Classification with SVM

```
1 # 1. Load and clean the dataset
2 import seaborn as sns
3 from sklearn.preprocessing import LabelEncoder
4 from sklearn.model_selection import train_test_split
5 from sklearn.preprocessing import StandardScaler
6 from sklearn.svm import SVC
7 import matplotlib.pyplot as plt
8
```

```


9 penguins = sns.load_dataset("penguins")          # Load Palmer Penguins
10 penguins = penguins.dropna()                    # Remove rows with missing
    values
11
12 # 2. Select features and encode labels
13 X = penguins[['bill_length_mm', 'flipper_length_mm']]
14 y = LabelEncoder().fit_transform(penguins['species'])
15 # species      integers 0, 1, 2
16
17 # 3. Split into training and test sets
18 X_train, X_test, y_train, y_test = train_test_split(
19     X, y, test_size=0.3, random_state=42
20 )
21
22 # 4. Scale features to zero mean & unit variance
23 scaler = StandardScaler().fit(X_train)
24 X_train_scaled = scaler.transform(X_train)
25
26 # 5. Train an R B F kernel SVM
27 #     gamma='scale'      1 / (n_features * var(X_train))
28 model = SVC(kernel='rbf', gamma='scale', C=1.0)
29 model.fit(X_train_scaled, y_train)
30
31 # 6. Visualize support vectors on training data
32 plt.figure(figsize=(6,6))
33 # Plot all training points (faded by species)
34 plt.scatter(X_train_scaled[:,0], X_train_scaled[:,1],
35             c=y_train, cmap='viridis', alpha=0.5, edgecolors='k')
36
37 # Overlay support vectors with bold circles
38 plt.scatter(model.support_vectors_[:,0], model.support_vectors_[:,1],
39             s=120, facecolors='none',
40             edgecolors='red', linewidths=2,
41             label='Support Vectors')
42
43 plt.xlabel('Bill Length (standardized)')
44 plt.ylabel('Flipper Length (standardized)')
45 plt.title('SVM Support Vectors on Penguin Features')
46 plt.legend()
47 plt.grid(True)
48 plt.show()

```

## Key Differences from Iris Dataset

- **Data Cleaning:** We drop missing values with `.dropna()`, ensuring no invalid rows are used.
- **Categorical Encoding:** Penguin species are converted from strings to integers (0, 1, 2) via `LabelEncoder`.
- **Feature Scaling:** Standardization ( $\mu = 0$ ,  $\sigma = 1$ ) is critical for SVM to treat each feature equally.
- **Gamma = 'scale':** Automatically sets  $\gamma = 1/(n\_features \times \text{Var}(X))$ , a sensible default for most problems.

- **Support Vectors:** These are the critical training points that lie exactly on the margin boundaries and define the classifier.



SVM decision boundaries and support vectors for penguin species classification

### Interpreting Support Vectors

- **Position:** Always lie on or between the two margin boundaries.
- **Significance:** Removing any single support vector will alter the optimal hyperplane.
- **Visual Clue:** In our plot, support vectors are marked by bold red circles.

## 5 Expert-Level Tips

### 5.1 When to Use the RBF Kernel

- **Pros:**
  - Can carve out highly complex, non-linear decision boundaries.
  - Automatically adapts to clusters of varying shapes and densities.
- **Cons:**
  - More computationally intensive than linear or low-degree polynomial kernels.
  - Requires careful tuning of  $\gamma$  and  $C$  to avoid overfitting or underfitting.
- **Sweet Spot:** Medium-sized datasets (up to 10,000 samples) where clear but non-linear clusters exist, and training time remains practical.

## 5.2 Gamma in Practice

Gamma Value	Decision Boundary	Risk / Behavior
$\gamma > 1$	Highly wiggly, tight around each point	Overfit: memorizes noise and outliers
Optimal ( $\approx$ “scale” or tuned)	Smooth, well-curved margins	Good trade-off: generalizes to new data
$\gamma < 0.01$	Almost linear, broad margins	Underfit: misses subtle patterns

## 5.3 Debugging SVM Performance

1. **Always scale features.** Use `StandardScaler` or `MinMaxScaler` so all dimensions contribute equally.
2. **Begin with `gamma='scale'`.** This default adapts  $\gamma$  to  $\frac{1}{n_{\text{features}} \times \text{Var}(X)}$ .
3. **If overfitting:**
  - Decrease  $\gamma$  to smooth the boundary.
  - Increase  $C$  to allow a wider margin (more slack).
4. **If underfitting:**
  - Increase  $\gamma$  for tighter local decision regions.
  - Consider a polynomial kernel (`kernel='poly'`) with degree 2–3.
5. **Monitor support vectors.** Too many support vectors often indicate a complex (potentially overfit) model.

## 6 Key Takeaways

### SVM Essentials

- **Kernel Choice:** `rbf` for complex patterns; `linear` for simple, high-dimensional data.
- **Gamma Wisdom:** Start with `gamma='scale'` or automated heuristics before manual tuning.
- **Visual Reality:** RBF creates “warped” boundaries that can be hard to interpret—always validate on a hold-out set.
- **Data Preparation:** Handle missing values, encode categorical labels, and standardize all features.

*“SVM with RBF is like teaching AI origami - it folds feature space to separate classes!”*