

Editorial: Assignment-W6A2: Core Concepts in Neural Networks and K-Means Clustering

Question 1 (MSQ)

A neural network is used to classify handwritten digits from the MNIST dataset.

Question: (Select all that apply)

Which of the following statements about the input representation for an FNN in this task are correct?

- A) The image is used as a 2D array of pixels
- B) The image is flattened into a 1D vector
- C) Each input neuron represents one pixel's brightness
- D) The input is a sequence of pixel values over time

Correct Answer: B), C)

Explanation:

The 28x28 image is flattened into a 784-length vector, with each input neuron representing a pixel's brightness.

Question 2 (MCQ)

You are working in a startup and are building a neural network for prediction. You observed that sometimes simpler neural network models were outperforming deeper ones.

Question: What is a key reason that simpler neural network models can sometimes outperform deeper ones? A) They always have higher accuracy

- B) They require more computation
- C) They are less likely to overfit with limited data
- D) They ignore regularization

Correct Answer: C)

Explanation:

Simpler models often generalize better, especially when data is limited, as deeper models may overfit without enough data or regularization

Question 3 (MCQ)

A research team at an AI company is training a neural network for a regression task. During backpropagation, they calculate the gradient of the loss with respect to the output as 0.4 and the gradient of the output with respect to a weight as 0.3.

Question:

What is the gradient of the loss with respect to the weight?

- A) 1.33
- B) 0.75
- C) 0.4
- D) 0.12

Correct Answer: D) 0.12

Explanation:

By the chain rule, the gradient of the loss with respect to the weight is the product of the two given gradients:

$$\partial L / \partial w = (\partial L / \partial y) \times (\partial y / \partial w) = 0.4 \times 0.3 = 0.12$$

Question 4 (MSQ)

A data scientist at a fruit-sorting company is designing a neural network to classify fruit images captured by a conveyor belt camera. The camera outputs RGB images where pixel values range from 0 to 255. Before feeding this data to the model, the scientist normalizes the pixel values to be in the range [0, 1].

Question:

Why is normalization of pixel values important when training a neural network on image data?

Select all that apply.

- A) It helps ensure all features contribute equally during learning
- B) It improves model convergence during gradient descent
- C) It removes background noise from the images
- D) It increases the resolution of the image data

Correct Answer:

☒ A), B)

Explanation:

Normalization ensures consistent scale across input features, which avoids one feature dominating the learning process. It also speeds up convergence by improving gradient behavior. It does **not** remove noise or change the image resolution.

Question 5 (MSQ)

You were assigned to modify an existing feedforward neural network (FNN) which is trained on the MNIST dataset. You were wondering...

Question: (Select all that apply)

What is the role of the hidden layers in the FNN?

- A) They directly output the class prediction
- B) They learn some pattern like basic shapes and digit parts
- C) They combine features into higher-level representations
- D) They only store the input data

Correct Answer: B), C)

Explanation:

Hidden layers in FNNs learn to detect some pattern like shapes and combine them into higher-level representations like digit parts, which are then used for classification.

Question 6 (MCQ)

A data scientist is designing a feedforward neural network to classify handwritten digits and must choose an activation function for the hidden layers.

Question:

Which activation function is used in the hidden layers of the FNN for MNIST, and why?

- A) Sigmoid, for non-linearity
- B) ReLU, for efficient non-linearity and sparse activations
- C) Softmax, to output probabilities
- D) Tanh, to center activations

Correct Answer: B)

Explanation:

ReLU is used because it efficiently introduces non-linearity and allows for sparse, fast computations by zeroing out negative values.

Question 7 (MSQ)

A team is developing a Recurrent Neural Network (RNN) to predict the next word in a sentence based on previous words.

Question: (Select all that apply)

In the context of RNNs, what does the hidden state represent?

- A) The current input only
- B) The memory of the network at a given time step
- C) Accumulated information from previous time steps
- D) The output probabilities

Correct Answer: B), C)

Explanation:

The hidden state in an RNN acts as memory, carrying information from previous time steps and integrating it with the current input¹.

Question 8 (MSQ)

A data scientist wants to segment customers into groups based on their purchasing behavior using K-means clustering.

Question: (Select all that apply)

Which of the following statements about K-means clustering are correct?

- A) It is a supervised learning algorithm
- B) It assigns each data point to the nearest centroid

- C) The number of clusters (K) must be specified in advance
- D) It guarantees finding the global minimum of the objective function

Correct Answer: B), C)

Explanation:

K-means is an unsupervised algorithm that assigns each data point to the nearest centroid, and the number of clusters K must be set before running the algorithm. It does not guarantee a global minimum due to random initialization.

Question 9 (MCQ)

A marketing analyst applies K-means clustering to customer data and notices that the resulting clusters are different each time the algorithm is run.

Question:

What is the most likely cause?

- A) The data is not normalized
- B) The number of clusters is too large
- C) The dataset is too small
- D) The algorithm uses random initialization of centroids

Correct Answer: D)

Explanation:

K-means uses random initialization for centroids, which can lead to different results on each run.

Question 10 (MSQ)

You are clustering data points using K-means. After several iterations, the assignment of points to clusters no longer changes.

Question: (Select all that apply)

What does this indicate?

- A) The algorithm has converged
- B) The centroids are still moving
- C) The algorithm has reached a local minimum.
- D) Further iterations will not change the clusters

Correct Answer: A), D)

Explanation:

When assignments stop changing, K-means has converged and further iterations will not alter the clusters.

Question 11 (MCQ)

A data scientist at a retail company has applied K-means clustering to group customer profiles based on their purchasing behavior. She now wants to assess how well the customers are grouped internally within each cluster.

Question:

Which metric should she use to evaluate how tightly grouped the customers are within each cluster?

- A) Inter-cluster distance
- B) Between-cluster variance
- C) Number of clusters
- D) Within-cluster sum of squares (WCSS)

Correct Answer: D)

Explanation:

K-means minimizes the within-cluster sum of squares, which measures how closely related the data points are within each cluster — a key metric for evaluating clustering quality.

Question 12 (MSQ)

You were working on a dataset with many outliers. You apply K-means clustering.

Question: (Select all that apply)

How might outliers affect the clustering result?

- A) Outliers can distort the position of centroids
- B) K-means is robust to outliers
- C) Outliers may form their own clusters
- D) The algorithm may converge more slowly

Correct Answer: A), C), D)

Explanation:

Outliers can pull centroids away from dense regions, may form their own clusters, and slow convergence. K-means is not robust to outliers.

Question 13 (MSQ)

A data analyst is using the elbow method to decide the number of clusters for K-means on a sales dataset.

Question: (Select all that apply)

What does the "elbow" point on the plot indicate?

- A) The point where adding more clusters yields diminishing returns
- B) The lowest inertia value
- C) The optimal number of clusters
- D) The maximum number of clusters possible

Correct Answer: A), C)

Explanation:

The elbow point shows where increasing K further does not significantly reduce inertia, suggesting the optimal K.

Question 14 (MSQ)

A machine learning engineer is reviewing the limitations of K-means clustering before applying it to a new dataset.

Question: (Select all that apply)

Which of the following are limitations of the standard K-means algorithm?

- A) Sensitive to initial centroid positions
- B) Can handle non-spherical clusters well
- C) Sensitive to outliers
- D) Requires specification of K in advance

Correct Answer: A), C), D)

Explanation:

K-means is sensitive to initial centroids and outliers, and the number of clusters K must be specified.

Question 15 (MSQ)

A researcher is considering K-means clustering for a dataset with mixed data types and several outliers.

Question: (Select all that apply)

Which situations are not suitable for K-means clustering? A) Data with non-spherical clusters

- B) Data with many outliers
- C) Data with categorical variables only
- D) Data that is well-separated and spherical

Correct Answer: A), B), C)

Explanation:

K-means works best for numerical, spherical, well-separated clusters and struggles with non-spherical clusters, outliers, and categorical data.