

AI Essentials: From Code Blueprints to Data Stories

*Object-Oriented Programming Principles, Real-World Visualizations, and Data Insights:
A Comprehensive Revision*

Minor In AI, IIT Ropar

20th June 2025

Okay, aspiring AI enthusiasts! Welcome to "Object-Oriented Programming Principles, Real-World Visualizations, and Data Insights: A Comprehensive Revision." This book is designed to take you on a friendly journey through some fundamental concepts that are crucial as you venture into the world of Artificial Intelligence. We'll connect these ideas to things you already know from everyday life and see how they empower us to build more organized code and understand complex data.

Think of this as a revision session, but presented in a way that builds intuition rather than just memorizing syntax. Let's dive in!

Chapter: Building Blocks and Seeing the Story in Data

Welcome, everyone! As you embark on your AI journey, you'll find that two core skills become incredibly valuable: knowing how to structure your code effectively and knowing how to understand and present the information hidden within your data. Today, we're going to revisit some foundational ideas in Object-Oriented Programming (OOP) and explore how we can use visualizations to uncover data insights.

Part 1: Bringing the Real World into Code - Object-Oriented Programming (OOP)

Why do we even need different ways of programming? You might have heard of functional programming (thinking like math problems), or procedural programming (thinking like a recipe with steps). Then came Object-Oriented Programming. The brilliant minds who invented it looked at the world around us and thought, "Hey, the real world has objects! Objects have properties and they can *do* things (operations). They relate to each other. Can we make our code work like that?"

Think about it:

- **Privacy:** We all have a public life (what we share on social media) and a private life (things known only to close friends or just ourselves). We build walls around our homes for privacy.
- **Inheritance:** We inherit things from our parents or ancestors – maybe money, property, or even physical traits and qualities (genes).
- **Relationships:** We have relationships with others – family bonds, friendships, professional connections.

OOP was designed to bring these real-world concepts – grouping properties and actions together, hiding certain details, inheriting traits – into programming languages. Why? Because we're programming things that interact with the real world!

While humans can understand context and even mistakes (like if I accidentally say "functional" when I mean "object-oriented," you probably know what I intended), computers are not so smart. They need strict rules and syntax. So, the challenge became: how do we translate these intuitive human concepts into a rigid language that a machine can understand?

This is where the principles of OOP come in. Let's explore them using a simple example: our beloved pets!

The Blueprint: Classes and Objects

Imagine you want to represent different pets in your program – a dog, a cat, a rabbit. They all have some common characteristics (like a name) and things they *do* (like make a sound). In OOP, we create a **Class** as a blueprint for these objects.

```
1 # We start by defining a class using the 'class' keyword (often shown in a
  # special color!)
2 class Pet:
3     # This is where we'll define the common characteristics and behaviors of
  # pets
4     pass # Placeholder for now
```

A **class** is just the definition, the mold. An **object** is an actual instance created from that mold – your specific dog named "Rocky", your cat named "Shiny".

Bringing an Object to Life: The `__init__` Method

When you get a new pet, the first thing you often do is name it! You also know it will have a certain temperament or 'mood' by default (maybe it's generally playful, sleepy, or cautious). This initial setup, or **initialization**, is handled by a special method in OOP called the **constructor**. In Python, this constructor method is specifically named `__init__` (that's *double* underscore, *init*, *double* underscore).

When you create a new object (an instance of the class), the `__init__` method is automatically called to set up its initial state. Think of buying an electronic device – it comes with some default settings already in place. `__init__` does this for your objects.

```
1 class Pet:
2     # The __init__ method is automatically called when you create a new Pet
  # object
3     def __init__(self, name): # 'self' refers to the specific object being
  # created
4         # Assign the name provided when creating the object
5         self.name = name
6         # Set a default mood (let's make it private for now!)
7         self.__mood = "sleepy"
```

Notice the `self` keyword. This is how we refer to the specific object we are currently working with. When you create `my_dog = Pet("Buddy")`, `self` inside `__init__` refers to `my_dog`. `self.name` means `my_dog`'s name.

Hiding Details: Public, Protected, and Private

Just like we have different levels of privacy in our lives, OOP languages allow us to control the visibility and accessibility of an object's properties (variables) and behaviors (methods).

- **Public:** By default, variables and methods are public. Like your name or what you post on social media, everyone can see and access them directly. In our `Pet` class, `self.name` is public.
- **Protected:** In Python, a single underscore prefix (e.g., `_mood`) is a *convention* indicating something is "protected". It's a hint to other programmers: "Hey, this is kind of private, maybe don't mess with it directly unless you know what you're doing." But technically, Python doesn't strictly prevent access. Think of it as telling a friend a secret but knowing they *could* tell others.
- **Private:** A double underscore prefix (e.g., `__mood`) signifies "private" in Python. This is a stronger signal. Python actually "mangles" the name internally (changes it to something

like `_ClassName__variableName`) to make it harder to access directly from outside the class. You *can* still access it using the mangled name, but the intent is clear: this is internal business! Think of it as a secret you absolutely do *not* want shared, walls built high.

Let's see this in action:

```
1 # Create a pet object
2 my_pet = Pet("Rocky")
3
4 # Accessing a public variable - Works fine!
5 print(my_pet.name)
6
7 # Attempting to access a private variable directly - Will likely cause an error
8 # !
9 # print(my_pet.__mood) # This line will give an AttributeError
```

Running this will show "Rocky" but fail on the `__mood` line because Python has mangled the name.

To access or change private variables, you typically create **getter** and **setter** methods within the class:

```
1 class Pet:
2     def __init__(self, name):
3         self.name = name
4         self.__mood = "sleepy" # Still private
5
6     # A getter method to access the private mood
7     def get_mood(self):
8         return self.__mood
9
10    # A method to change the mood (implicitly acting as a setter here)
11    def play(self):
12        self.__mood = "excited"
13        print(f"{self.name} is now {self.__mood}!")
```

Now, you access the mood indirectly:

```
1 my_pet = Pet("Rocky")
2 print(f"Initial mood: {my_pet.get_mood()}")
3 my_pet.play()
4 print(f"Mood after playing: {my_pet.get_mood()}")
```

This controlled access (**Encapsulation**) is a key OOP principle – bundling data (`__mood`) and the methods that operate on it (`get_mood`, `play`) within a single unit (the class) and controlling outside access.

Building on What Exists: Inheritance

Pets are diverse! We have dogs, cats, rabbits, etc. A `Dog` is a type of `Pet`. It inherits all the general properties of a `Pet` (like having a name and a mood), but it also has its own specific properties (like breed) and behaviors (like barking).

This is **Inheritance**. A new class (**derived class** or **child class**) can inherit attributes and methods from an existing class (**base class** or **parent class**).

```
1 # Dog is a type of Pet, so it inherits from Pet
2 class Dog(Pet):
3     # The Dog class also has its own __init__ method
4     def __init__(self, name, breed):
5         # But first, we need to call the parent's __init__ to set up inherited
6         # properties!
```

```

6         # The 'super()' function refers to the parent class
7         super().__init__(name) # Calls Pet's __init__ to set the name and
        default mood
8
9         # Now add properties specific to Dog
10        self.breed = breed
11
12        # Add a behavior specific to Dog
13        def bark(self):
14            print(f"{self.name} the {self.breed} says Bow bow!")
15
16        # Override or add a method to show mood (uses the inherited get_mood
        implicitly)
17        def show_mood(self):
18            # We can access inherited methods using self
19            print(f"{self.name} is feeling right now {self.get_mood()}")

```

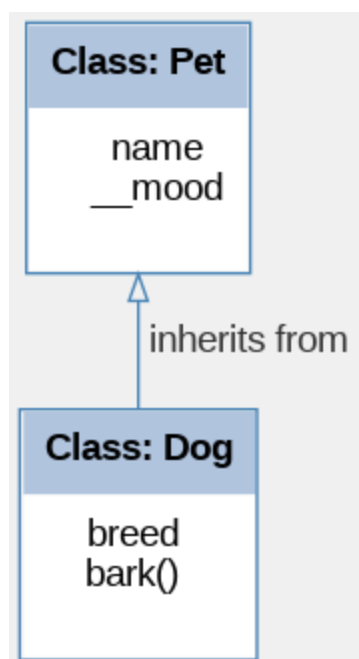


Figure 1: Diagram showing "Class: Pet" at the top, with an arrow labeled "inherits from" pointing down to "Class: Dog". The Dog box contains "breed" and "bark()" which are unique to it.

When we create a Dog object, its `__init__` is called. Inside, `super().__init__(name)` ensures the Pet part of the Dog object is initialized correctly. Then, the Dog's `__init__` adds the `breed`.

```

1 # Create a Dog object
2 my_dog = Dog("Ronick", "Dalmatian")
3
4 # Use inherited methods (from Pet)
5 print(my_dog.name) # Inherited attribute
6 my_dog.play() # Inherited method (modified mood)
7
8 # Use Dog-specific methods
9 my_dog.bark()
10 my_dog.show_mood() # Uses inherited get_mood internally

```

This covers the core OOP concepts discussed: Classes, Objects, `__init__` (Constructor), `self`, Public/Protected/Private members, Mangling, Encapsulation, Inheritance, and `super()`.

Part 2: Seeing the Story in Data - Visualizations

Once you have your data, the next crucial step, especially in AI, is to understand it. Numbers alone can be overwhelming. Visualizations turn data into pictures, making patterns, trends, and insights jump out.

We'll quickly revisit some basic plots and then explore some more powerful ones offered by libraries like Matplotlib (often imported as `plt`) and Seaborn (often imported as `sns`).

Let's start with a simple dataset loaded into a Pandas DataFrame:

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 import pandas as pd
4
5 # Sample data (replace with actual data loading in a real scenario)
6 data = {'Category': ['A', 'B', 'C', 'D', 'E'],
7         'Values': [10, 25, 15, 30, 20],
8         'Numerical1': [150, 160, 165, 170, 180],
9         'Numerical2': [50, 65, 70, 75, 80]}
10 df_basic = pd.DataFrame(data)
11
12 # Basic plots (using plt) - Quick review of names!
13 # plt.plot(df_basic['Category'], df_basic['Values']) # Line Plot
14 # plt.scatter(df_basic['Numerical1'], df_basic['Numerical2']) # Scatter Plot
15 # plt.bar(df_basic['Category'], df_basic['Values']) # Bar Plot
16 # plt.hist(df_basic['Values']) # Histogram
17 # plt.pie(df_basic['Values'], labels=df_basic['Category']) # Pie Chart
18 # plt.boxplot(df_basic['Values']) # Box Plot
19 # plt.show() # Display the plots
```

These basic plots are powerful for showing relationships between two variables or the distribution of a single variable. But what if you have many variables?

Seeing All Pairwise Relationships: The Pair Plot

Imagine you have data about students, including their height, weight, age, and exam scores. You want to see how each of these relates to every other one. Seaborn's `pairplot` comes to the rescue. It automatically generates a grid of plots showing the relationship between each pair of numerical columns in your DataFrame.

```
1 # Sample student data
2 data_students = {'Height': [150, 160, 165, 170, 180, 155, 175],
3                  'Weight': [50, 65, 70, 75, 80, 55, 78],
4                  'Age': [25, 30, 35, 40, 45, 28, 38],
5                  'Score': [60, 75, 80, 85, 90, 68, 88]}
6 df_students = pd.DataFrame(data_students)
7
8 # Create a pair plot
9 # sns.pairplot(data=df_students)
10 # plt.show() # Remember to show the plot!
```

The plots *off* the diagonal show scatter plots for each pair of different variables (e.g., Height vs. Weight). The *diagonal* plots show the **distribution** of each single variable, typically as histograms or KDEs.

Comparing Across Categories: The Box Plot and Swarm Plot

Often, you want to compare a numerical variable *within* each category.

```
1 # Sample gym membership data
2 data_gym = {'Membership': ['Basic', 'Premium', 'Basic', 'Premium', 'Basic', 'Premium', 'Basic', 'Premium'],
3             'Gym': [1, 2, 1, 2, 1, 2, 1, 2]}
```

```

3         'Satisfaction': [3.0, 4.5, 3.2, 4.8, 2.8, 4.7, 3.1, 4.6],
4         'Calories Burnt': [200, 500, 250, 550, 230, 530, 260, 560]}
5 df_gym = pd.DataFrame(data_gym)

```

Box Plot (sns.boxplot): This is excellent for comparing the *summary statistics* of a numerical distribution across different categories.

```

1 # Plot Satisfaction vs Membership using a Box Plot
2 sns.boxplot(x='Membership', y='Satisfaction', data=df_gym)
3 plt.title('Satisfaction Rating by Membership Type')
4 # plt.show()

```

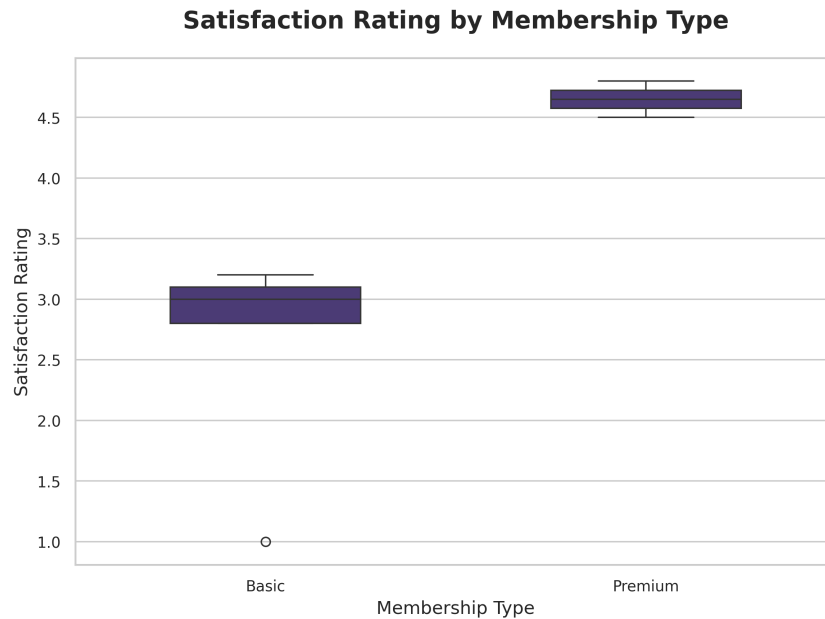


Figure 2: A box plot showing two boxes side-by-side, one for 'Basic' and one for 'Premium' membership satisfaction ratings. The plot shows median lines, boxes (IQR), and whiskers.

How to read a box plot:

- The box is the **Interquartile Range (IQR)** – the middle 50% of data.
- The line inside is the **Median**.
- The "whiskers" show the rest of the data's range.
- Points outside the whiskers are potential **Outliers**.

Swarm Plot (sns.swarmplot): While box plots summarize, swarm plots show *every single data point*. This is great for smaller datasets where you want to see the density of individual points.

```

1 # Sample shopping data
2 data_shop = {'Time of Day': ['Morning', 'Afternoon', 'Evening', 'Morning', '
        Afternoon', 'Evening', 'Afternoon', 'Morning'],
3              'Amount Spent': [20, 150, 80, 30, 200, 90, 180, 25]}
4 df_shop = pd.DataFrame(data_shop)
5
6 # Plot Amount Spent vs Time of Day using a Swarm Plot
7 # sns.swarmplot(x='Time of Day', y='Amount Spent', data=df_shop)
8 # plt.title('Amount Spent at Different Times of Day')
9 # plt.show()

```

Finding Connections: Correlation and Heatmaps

To quickly see how strongly numerical features are related, we use **correlation**. A value near +1 implies a strong positive linear relationship, near -1 implies a strong negative one, and near 0 implies a weak one. You can calculate this with the `.corr()` method.

```
1 # Sample hobby data (all numerical)
2 data_hobby = {'Introversion Score': [7, 5, 8, 6, 4],
3               'Gaming Hours/Week': [10, 8, 12, 9, 7],
4               'Reading Hours/Week': [3, 5, 2, 4, 6]}
5 df_hobby = pd.DataFrame(data_hobby)
6
7 # Calculate the correlation matrix
8 correlation_matrix = df_hobby.corr()
9 print(correlation_matrix)
```

Looking at a table of numbers can be hard. A **Heatmap** visualizes this matrix using colors.

```
1 # Plot the correlation matrix using a heatmap
2 sns.heatmap(data=correlation_matrix, annot=True, cmap='coolwarm')
3 plt.title('Correlation Heatmap of Hobby Data')
4 # plt.show()
```

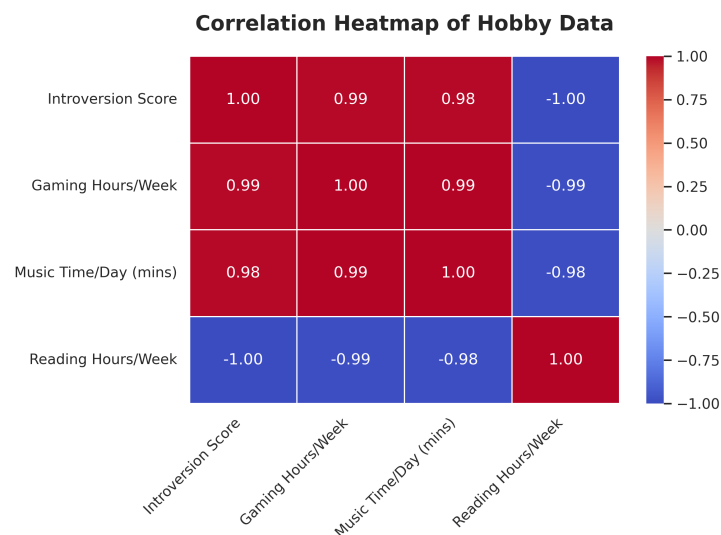


Figure 3: A heatmap of the correlation matrix. Colors represent the correlation values. Cells with high positive correlation (close to 1) and high negative correlation (close to -1) should stand out visually.

Self-reflection moment

Remember that poll about Pearson correlation? The Pearson coefficient is just a specific way to calculate this correlation! If the value is small (close to 0), the linear relationship is weak. The sign (+ or -) tells you the direction. If a -0.1 relationship is compared to a +0.2, both are weak, but +0.2 is slightly stronger than -0.1 in terms of magnitude.

Connecting the Dots: Upcoming Project

We've touched upon how OOP helps structure your code into logical, reusable units (like our `Pet` and `Dog` classes) and how visualizations help you extract meaningful insights from data (seeing relationships, distributions, and comparisons).

As we move forward, especially starting next week, we'll bring these concepts together. We'll take a real dataset, use our programming skills, analyze it using various techniques, and visualize our findings to build a simple mini-project. This will solidify how the foundational concepts you've learned in modules A and B are applied in practice.

Conclusion

Today, we took a step back to appreciate *why* we use OOP, seeing how it mirrors aspects of the real world like privacy and inheritance to create better-structured code. We revisited classes, objects, constructors (`__init__`), the self-reference, and controlling access with public, protected, and private members.

Then, we zoomed out to look at data itself, understanding how different plots – from basic line and scatter plots to more advanced pair plots, box plots, swarm plots, and heatmaps – tell unique stories about our data. We saw how to choose the right visualization based on whether we're looking at relationships between many variables, comparing distributions across categories, or identifying correlations.

These skills are not just for theory; they are your essential tools for building, understanding, and explaining AI models. Keep practicing, keep exploring the different parameters these plotting functions offer, and most importantly, keep asking "what does this mean?" when you look at your code or your data.

Keep hydrating, get some physical activity in, and prepare for some exciting practical application next week!

Good night, and happy learning!