



Post-Session Notes: Revision – Evaluating Metrics and Result Analysis, Testing and Training



1. Project Introduction & Evaluation Context

- The project focuses on analyzing healthcare-related data to uncover patterns and model relationships in variables like:
 - Hours of sleep
 - Sleep quality
 - Stress level
 - Daily steps
 - Calorie intake
 - This session emphasized the importance of interpreting model results, handling class imbalance, and evaluating performance with appropriate metrics.
 - Earlier steps like data preprocessing, visualization, and basic modeling (SVM, Logistic Regression) laid the groundwork.
-



2. Clustering Techniques: K-Means & DBSCAN

- ◆ K-Means Clustering

- Goal: Group data into clusters based on distance to centroids (means).

- Key parameters:

- **n_clusters**: Predefined number of clusters.
- **random_state**: For reproducibility.

- Application:

- Used on features like sleep hours, sleep quality, and stress level.
- Helped in identifying behavioral patterns among users.

- Evaluation with Silhouette Score:

- Range: -1 to 1
- >0.5 = good clustering
- ~ 0.3 = weak but acceptable
- Clusters showed overlap and scattering, indicating complex relationships not fully captured by K-means.

◆ DBSCAN Clustering

- Goal: Density-based clustering that doesn't require predefined clusters.

- Key Parameters:

- **eps**: Maximum distance to consider for neighborhood.

- **min_samples**: Minimum points to form a cluster.
 - **Advantages:**
 - Handles noise/outliers explicitly.
 - Finds non-linear, irregularly shaped clusters.
 - **Application:**
 - Applied on calorie intake and daily steps.
 - Resulted in clusters + noise points.
 - Tuning **eps** drastically affected number and quality of clusters.
-

3. Evaluation Metrics: Going Beyond Accuracy

- ◆ **Why Accuracy Isn't Enough**
 - In imbalanced datasets, high accuracy can be misleading.
 - E.g., 98% accuracy may hide the fact that the minority class is barely predicted.
- ◆ **Key Metrics:**
 - **Precision:** How many predicted positives are actually correct?
(*Focus: False Positives*)

- **Recall:** How many actual positives are correctly predicted?
(*Focus: False Negatives*)
- **F1 Score:** Harmonic mean of precision and recall – balances both.
- **Confusion Matrix:** Breaks down predictions into:
 - True Positives
 - False Positives
 - True Negatives
 - False Negatives

Used a logistic regression model to demonstrate how a model could score high accuracy but low precision/recall for minority classes.

4. Handling Class Imbalance

- Imbalanced data (e.g., very few samples in one class) affects model performance and metric fairness.
- **Solutions:**
 - Use `class_weight='balanced'` in models (like Logistic Regression) to penalize majority class overconfidence.
 - Consider oversampling or undersampling (not covered deeply here but conceptually important).

```
model = LogisticRegression(class_weight='balanced')
model.fit(X_train, y_train)
```

5. Train-Test Split vs Cross-Validation

♦ Train-Test Split

- Simple partition of data into training and testing subsets.
- Can lead to biased results depending on the split.

♦ Cross-Validation

- K-Fold CV splits data into multiple subsets.
- Each fold is used once as a test set, others as training.
- Reduces overfitting risk and provides average metric performance.

```
from sklearn.model_selection import cross_val_score
scores = cross_val_score(model, X, y, cv=5, scoring='f1')
print("Mean F1 score:", scores.mean())
```

- Epoch vs Cross-Validation:
 - Epoch: One full pass through training data (used in deep learning).
 - Cross-validation: Evaluation technique to test model robustness.
-

6. Interpreting Clustering Output

- Cluster results from both KMeans and DBSCAN require domain interpretation:
 - What does each cluster represent?
 - How does it help in understanding user behavior?
 - Challenges:
 - Overlapping clusters
 - Arbitrary shapes in data
 - Defining meaningful cluster labels
-

7. Summary of Key Coding Techniques

```
from sklearn.cluster import KMeans, DBSCAN
from sklearn.metrics import silhouette_score, confusion_matrix, f1_score
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LogisticRegression
```

KMeans Clustering

```
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans.fit(X)
labels = kmeans.labels_
silhouette = silhouette_score(X, labels)
```

DBSCAN Clustering

```
dbscan = DBSCAN(eps=0.5, min_samples=20)
```

```
dbscan_clusters = dbscan.fit_predict(X)
```

```
# Logistic Regression with Class Weights
```

```
logreg = LogisticRegression(class_weight='balanced')
```

```
logreg.fit(X_train, y_train)
```

```
# Cross-validation
```

```
f1_scores = cross_val_score(logreg, X, y, cv=5, scoring='f1')
```

8. Key Takeaways & Best Practices

- Never rely solely on accuracy, especially with imbalanced data.
 - Always look at precision, recall, and F1 for classification models.
 - Silhouette score is an effective way to evaluate clustering quality.
 - Experiment with clustering parameters and features.
 - DBSCAN is powerful for noisy, irregular data.
 - Cross-validation offers reliable performance assessment.
 - Model evaluation is not just technical — it needs contextual understanding.
-

Conclusion

This session integrated clustering techniques, evaluation metrics, and testing strategies into the broader healthcare dataset project. It emphasized

robustness, interpretability, and thoughtful metric selection—all vital in building real-world, dependable ML models.
