

Unlocking Image Data - Pixels, Colors, and Our First Steps

Minor In AI, IIT Ropar

15th May, 2025

Welcome, aspiring AI enthusiasts! Today, we embark on a visually exciting journey – diving into the world of images. You see images everywhere, on your phones, computers, in books, and all around you. But how do these images, full of vibrant colors and intricate details, get stored and understood by a computer? How does AI even begin to *see*?

Let's start with a simple exercise, a case study, much like we did in our class. Take a look at this image:



Figure 1: The Entrance to IIT Ropar

Most of us instantly recognize this, right? It's the entrance to IIT Ropar campus. Our human brains process this complex visual information effortlessly. We see buildings, pathways, trees, the sky, maybe even cars or people if they were there. But what does a computer see? How does it 'understand' this is a campus entrance?

Think about zooming into an image. What happens? If you zoom in too much, you start seeing tiny squares, right? These are the fundamental building blocks of any digital image: **pixels**.

1 The Digital Canvas: Pixels and Their Colors

Imagine your image is a giant grid, like graph paper. Each tiny square on this grid is a pixel. Every single pixel holds information about the color it should display. When you put millions of these colored pixels together in the right order, they form the complete image we see.

But what kind of color information does a pixel hold?

Most digital images use a standard way to represent color: the **RGB** model. This stands for **Red**, **Green**, and **Blue**. The idea is that almost any color we see can be created by mixing different amounts of these three primary colors of light.

So, each pixel isn't just "blue" or "green." Instead, it holds a recipe for its color, specifying how much Red, how much Green, and how much Blue to mix.

2 Peeking Inside Pixels: Binary and Intensity

How does a computer store these amounts of Red, Green, and Blue? This is where we get into the exciting world of how computers represent numbers. You might recall from previous discussions that computers fundamentally understand things in terms of zeros and ones – binary digits, or bits.

A common way to store the intensity (or amount) of each color (Red, Green, or Blue) for a pixel is using **8-bit encoding**. What does 8-bit mean? It means we use 8 binary digits (eight 0s or 1s) to represent a number.

Think back to how we convert binary to decimal numbers using powers of 2:

- $2^0 = 1$
- $2^1 = 2$
- $2^2 = 4$
- ... and so on.

If we have 8 bits, the smallest number we can represent is when all 8 bits are 0: 00000000. In decimal, this is **0**.

The largest number we can represent is when all 8 bits are 1: 11111111. If you add up the powers of 2 from 2^0 to 2^7 , you get $1 + 2 + 4 + 8 + 16 + 32 + 64 + 128$, which sums up to **255**.

So, with 8-bit encoding, we can represent values from **0 to 255**.

This range (0 to 255) is used for the intensity of each color channel (Red, Green, and Blue) in a pixel.

- A value of **0** for a color means *none* of that color is present.
- A value of **255** means the *maximum* amount of that color is present.

3 How RGB Values Combine

Now, consider a single pixel. It has three values: one for Red, one for Green, and one for Blue, each ranging from 0 to 255.

Let's look at an example: (128, 64, 200)

- The Red value is 128 (somewhere in the middle).
- The Green value is 64 (relatively low).
- The Blue value is 200 (quite high).

Based on these values, which color do you think is dominant? The highest value is for Blue (200). So, this pixel would likely have a color that leans towards blue – maybe a shade of purple or deep blue, depending on the exact mix. We might say this pixel is on "Team Blue"!

What about these?

- (0, 0, 0): No Red, no Green, no Blue. This results in **Black**. (Remember, 0 is minimum intensity).
- (255, 255, 255): Maximum Red, maximum Green, maximum Blue. This results in **White**. (Remember, 255 is maximum intensity).

- (255, 0, 0): Maximum Red, no Green, no Blue. This results in pure **Red**.
- (0, 255, 0): Pure **Green**.
- (0, 0, 255): Pure **Blue**.

By combining these values in different proportions, we can create over 16 million unique colors (256 possibilities for Red \times 256 for Green \times 256 for Blue \approx 16.7 million). This is often referred to as "True Color" or 24-bit color (8 bits for each of the three channels).

4 Loading and Analyzing Images with Code

Now that we understand pixels and RGB representation, let's see how we can use code to work with images. We'll use a powerful library called **OpenCV** (imported as `cv2` in Python), which is specifically designed for computer vision tasks.

First, we need to load an image and start inspecting its pixel data.

```

1 import cv2
2 import numpy as np # We'll use numpy for calculations like summing
3 import matplotlib.pyplot as plt # For displaying images and plots
4
5 # Load the image
6 # OpenCV loads images by default in BGR format (Blue, Green, Red)
7 image = cv2.imread('iit-ropar.jpg') # Replace 'iit-ropar.jpg' with your
   image file name
8
9 # --- IMPORTANT STEP ---
10 # Matplotlib and many other libraries expect images in RGB format.
11 # OpenCV loads in BGR. We need to convert it for correct display and
   processing
12 image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
13
14 # Now, 'image_rgb' is a NumPy array representing the image data
15 # Its shape is (height, width, channels) where channels are (Red, Green
   , Blue)
16
17 # Let's look at the image and a simple analysis: total color intensity
18 # Display the original image (using the RGB converted version for
   correct colors)
19 plt.figure()
20 plt.imshow(image_rgb) # Use image_rgb for correct colors
21 plt.title("Original RGB Image")
22 plt.axis('off') # Turn off axes numbers/ticks
23 plt.show()
24
25 # Now, let's analyze the color intensity
26 # We can split the image array into its R, G, and B channels
27 # Remember the array structure is (rows, columns, channels)
28 # [:, :, 0] means "all rows, all columns, the 0th channel" (Red)
29 # [:, :, 1] means "all rows, all columns, the 1st channel" (Green)
30 # [:, :, 2] means "all rows, all columns, the 2nd channel" (Blue)
31 r = image_rgb[:, :, 0] # This is now a 2D array (height x width) of Red
   pixel values
32 g = image_rgb[:, :, 1] # 2D array of Green pixel values
33 b = image_rgb[:, :, 2] # 2D array of Blue pixel values
34
35 # Let's sum up the intensity values for each channel across the entire
   image

```

```

36 total_r = np.sum(r)
37 total_g = np.sum(g)
38 total_b = np.sum(b)
39
40 # Plot a bar chart to visualize the total intensity of each color
41 colors = ['Red', 'Green', 'Blue']
42 values = [total_r, total_g, total_b]
43
44 plt.figure(figsize=(6, 4))
45 plt.bar(colors, values, color=['red', 'green', 'blue'])
46 plt.title('Total Color Intensity in Image')
47 plt.ylabel('Sum of Pixel Values (0-255)')
48 plt.tight_layout() # Adjust layout to prevent labels overlapping
49 plt.show()

```

Listing 1: Loading an image and analyzing color intensities

Explanation of the Code:

1. We import the necessary libraries: `cv2` for image operations, `numpy` for numerical operations on the image data (which is stored as a NumPy array), and `matplotlib.pyplot` for plotting.
2. `cv2.imread('iit-ropar.jpg')` loads the image file. A crucial point here, and something that often trips up beginners, is that **OpenCV loads images in BGR order by default**, not RGB.
3. `cv2.cvtColor(image, cv2.COLOR_BGR2RGB)` is essential! Since `matplotlib` (which we use to display the image) and many standard image processing techniques assume RGB order, we *must* convert the image from OpenCV's BGR format to RGB. If you skip this, the colors will look wrong (blues will appear red, and vice versa).
4. The loaded image (`image_rgb`) is a NumPy array. Its `shape` tells you its dimensions: (`height`, `width`, `channels`). For an RGB image, `channels` will be 3.
5. We access the individual color channels using NumPy slicing: `[:, :, 0]` selects all rows, all columns, but only the *first* channel (which is Red in our converted `image_rgb`). Similarly, `[:, :, 1]` gets Green, and `[:, :, 2]` gets Blue. Each `r`, `g`, and `b` is now a 2D array representing the intensity of that color across the image.
6. `np.sum()` adds up all the pixel intensity values within each channel's 2D array. This gives us a single number representing the total "amount" of Red, Green, and Blue across the entire image.
7. Finally, we use `matplotlib` to display the correctly colored `image_rgb` and a bar chart showing the total intensity sums for Red, Green, and Blue.

Running this code on the IIT Ropar image, you'd likely see a bar chart where Green has the highest total intensity, followed by Blue, and then Red. This makes sense for a daytime outdoor shot with lots of trees, grass, and sky! If you loaded a different image, say one with a lot of water or night sky, you might see the Blue bar being the dominant one, as some of your classmates observed during the lecture.

5 Visualizing the Channels

To really understand how RGB works, let's visualize the individual color channels. We can display each channel as a grayscale image. A higher pixel value (closer to 255) in the channel will appear brighter in the grayscale image, and a lower value (closer to 0) will appear darker.

```

1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 # Load and convert to RGB (assuming the same image file)
6 image = cv2.imread('iit-ropar.jpg')
7 image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
8
9 # Display shape and data type (optional, but good to see)
10 print("Shape:", image_rgb.shape) # (height, width, channels)
11 print("Data Type:", image_rgb.dtype) # e.g., uint8 for 0-255 values
12
13 # Split RGB channels (again)
14 r, g, b = image_rgb[:, :, 0], image_rgb[:, :, 1], image_rgb[:, :, 2]
15
16 # Plot RGB channels separately as grayscale images
17 plt.figure(figsize=(10,3)) # Adjust figure size for 3 plots
18 for i, channel in enumerate([r, g, b]):
19     plt.subplot(1, 3, i+1) # Create 1 row, 3 columns of plots, select
20     # the i+1 plot
21     plt.imshow(channel, cmap='gray') # Use 'gray' colormap to show
22     # intensity
23     plt.title(['Red', 'Green', 'Blue'][i] + ' Channel (Grayscale)') #
24     # Set title
25     plt.axis('off') # Hide axes
26 plt.tight_layout()
27 plt.show()

```

Listing 2: Visualizing individual RGB channels as grayscale images

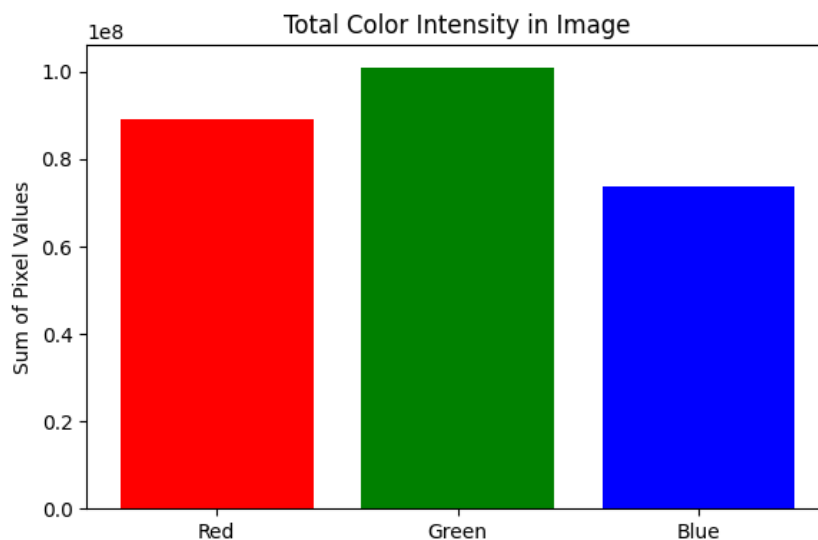


Figure 2: Visualizing Individual Color Channels as Grayscale Images

Looking at these grayscale images, you can see which parts of the original image had high intensity for each specific color. In the Red channel grayscale, areas with more red appear brighter. In the Green channel, areas with more green (like the grass and trees) are brightest. In the Blue channel, the sky area is likely the brightest. This visualization helps us understand the contribution of each color component to the final image.

6 Isolating the Color Contribution

We can take this a step further and create images that *only* show the contribution of a single color channel, while zeroing out the others. This shows how the intensity of one channel looks when only that color is allowed to be visible.

```
1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 # Load and convert to RGB
6 image = cv2.imread('iit-ropar.jpg')
7 image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
8
9 # Split channels
10 r, g, b = image_rgb[:, :, 0], image_rgb[:, :, 1], image_rgb[:, :, 2]
11
12 # Create a 'zero channel' - a 2D array of the same shape as r, g, or b,
13   # filled with zeros
14 zeros = np.zeros_like(r) # Or np.zeros(r.shape, dtype=r.dtype)
15
16 # Stack each original channel with two zeroed channels to isolate color
17 # np.stack joins arrays along a new axis. axis=2 means we stack along
18   # the channel dimension.
19 # For red_image, we stack (Red channel, Zero channel, Zero channel)
20 red_image = np.stack([r, zeros, zeros], axis=2)
21 # For green_image, we stack (Zero channel, Green channel, Zero channel)
22 green_image = np.stack([zeros, g, zeros], axis=2)
23 # For blue_image, we stack (Zero channel, Zero channel, Blue channel)
24 blue_image = np.stack([zeros, zeros, b], axis=2)
25
26 # Plot all three isolated color images
27 plt.figure(figsize=(10, 3))
28 for i, channel_img in enumerate([red_image, green_image, blue_image]):
29     plt.subplot(1, 3, i+1)
30     plt.imshow(channel_img) # Matplotlib will display this using RGB
31     # interpretation
32     plt.title(['Red', 'Green', 'Blue'][i] + ' Channel Contribution')
33     plt.axis('off')
34
35 plt.tight_layout()
36 plt.show()
```

Listing 3: Isolating color contributions of RGB channels



Figure 3: Visualizing Isolated Color Channel Contributions

Look at the Green Channel isolated image. You can clearly see how the green information contributes to the grass and trees. The sky might look quite dark, as it doesn't have much green

component. Similarly, in the Blue Channel image, the sky should appear prominently. This truly shows the power of combining these three channels to form the final colorful image!

7 Beyond Basic Colors

This understanding of pixels, RGB values, and channels is fundamental to almost everything we do in computer vision. Whether it's applying simple filters, detecting edges, or building complex deep learning models to recognize objects, we are constantly working with and manipulating these pixel values.

In the next chapter, we'll explore other ways images can be represented and start to look at techniques that analyze patterns and features within these pixel arrangements, moving us closer to how a computer can truly "see" and understand the world in images. Get ready, because working with images is visually intuitive and incredibly rewarding!

—