

Minor in AI

MLP => CNN

1 Introduction: The Birth of Digital Vision

Real-World Case Study

The Post Office Problem (1980s):

Imagine mountains of handwritten letters arriving daily at a post office. Workers had to manually read ZIP codes, a tedious and error-prone process. Computer scientist Yann LeCun faced this challenge: **“How can machines learn to read handwritten digits as accurately as humans?”**

- **Challenge:** Same digit (e.g., ‘5’) varies in shape, size, and tilt
- **Solution Path:** Create neural networks that learn patterns from examples
- **Breakthrough:** MNIST dataset - 70,000 labeled digit images

2 Multi-Layer Perceptrons (MLPs): Digital Brain Cells

2.1 Anatomy of an MLP

An MLP works like a team of detectives:

- **Input Layer:** 784 “sensors” (28x28 pixel values)
- **Hidden Layers:** Detect patterns (edges → curves → shapes)
- **Output Layer:** 10 “jurors” voting for digit 0-9

Why Matrix Math?

Neural networks = Fancy calculators:

Layers communicate through matrix operations. For layer 1 with 4 neurons:

$$\begin{bmatrix} w_{11} & w_{12} & \cdots & w_{1,784} \\ w_{21} & w_{22} & \cdots & w_{2,784} \\ w_{31} & w_{32} & \cdots & w_{3,784} \\ w_{41} & w_{42} & \cdots & w_{4,784} \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_{784} \end{bmatrix} = \begin{bmatrix} z_1 \\ z_2 \\ z_3 \\ z_4 \end{bmatrix}$$

Each weight w_{ij} is a **connection strength** between pixel j and neuron i .

2.2 Code Walkthrough: Matrix Operations

Listing 1: Python Implementation of Layer Calculation

```
import numpy as np

# Flatten 28x28 image into 784 numbers (0-255)
pixel_values = np.random.rand(784) # Simulated MNIST image

# First hidden layer weights (4 neurons)
```

```
weights_layer1 = np.random.randn(4, 784) * 0.1 # Small random values

# Calculate neuron activations
raw_output = np.dot(weights_layer1, pixel_values) # Matrix
multiplication
activated_output = 1 / (1 + np.exp(-raw_output)) # Sigmoid "squish"

print(f"Neuron activations: {activated_output}")
```

Code Explanation

What's happening here:

- `np.random.rand(784)`: Simulates normalized pixel values (0=black, 1=white)
- `np.random.randn(4, 784)`: Creates 4×784 weight matrix (Gaussian distribution)
- `np.dot()`: Matrix multiplication (input pixels × weights)
- `1/(1+np.exp(-x))`: Sigmoid activation (squishes values to 0-1)

3 Training Dynamics: Teaching Neural Networks

Three key ingredients:

1. Batches: The Classroom Analogy

Like studying flashcards in small sets rather than entire textbooks


- Process 64-256 images at once (mini-batches)
- Why? *Memory efficiency* and *frequent updates*
- Example: 60,000 MNIST images → 600 batches of 100 images
- Code analogy: `for batch in dataset: train(batch)`

2. Epochs: The Semester System

Complete passes through all training data

- 1 epoch = 1 full cycle through all 60,000 images
- Typical training: 10-100 epochs
- Why multiple? *Gradual learning* like semester revisions

Loss decreases with each epoch



3. Backpropagation: The Blame Game

"Which neuron made the biggest mistake?"

- Calculates error gradients using chain rule
- Flows backward: Output → Hidden layers → Input

Weight Update Equation

$$\underbrace{W_{new}}_{\text{Updated weights}} = \underbrace{W_{old}}_{\text{Current weights}} - \underbrace{\alpha}_{\text{Learning rate}} \times \underbrace{\frac{\partial L}{\partial W}}_{\text{Error gradient}}$$

Key components:

- **Learning Rate (α)**: Controls step size (0.001 typical)
 - **Too small**: Slow learning (turtle pace)
 - **Too big**: Overshooting (jumping valleys)
- **Gradient ($\frac{\partial L}{\partial W}$)**: Direction to reduce error

4 The Flattening Problem: Why MLPs Struggle

Spatial Relationships

The Jigsaw Puzzle Analogy: Imagine solving two different puzzles:

- **Intact Puzzle**: Edges match, colors flow naturally
- **Shuffled Pieces**: No spatial relationships, pure color matching

MLPs work with shuffled pieces! When we flatten a 28×28 image into 784 pixels:

- **Lost:**
 - Local patterns (horizontal edge at top, curve at bottom-right)
 - Relative positions (nose above mouth in face recognition)
 - Translation invariance (recognize '5' anywhere in image)
- **Kept:**
 - Global brightness (dark pixels vs light background)
 - Individual pixel intensities (0-255 values)

Why This Matters for Digits: A handwritten '8' requires:

- Top loop + bottom loop + middle cross
- Spatial arrangement crucial

MLP's Limited Perspective

"I see pixels, not patterns!"





An MLP neuron might learn:

- Pixel 123 is dark \rightarrow maybe a '7'
- Pixel 456 is bright \rightarrow maybe a '0'

But misses the **relationships** between pixels 123 and 456!

5 The Path to Convolutional Networks

Key Takeaways

-  **The Flattening Paradox:**
 - MLPs require 1D vectors \rightarrow Destroys 2D structure
 - Like reading a book by scrambling all words
-  **Training Machinery:**
 - Batches: 64-256 samples (memory-efficient learning)
 - Backpropagation: Error feedback through layers
 - Weight updates: $W = W - \alpha \nabla L$ (learning rate)
-  **Probability Conversion:**
 - Softmax: $p_i = \frac{e^{q_i}}{\sum e^{q_j}}$
 - Converts scores \rightarrow Probability distribution
-  **The Spatial Crisis:**
 - MLP accuracy plateau: 97-98% on MNIST
 - Real-world images need **hierarchical pattern recognition**

The CNN Revolution: Convolutional Neural Networks solve spatial blindness through:

- **Local receptive fields:** Detect edges/textures
- **Weight sharing:** Same pattern detectors across image
- **Pooling:** Spatial hierarchy preservation

Historical Insight: Yann LeCun's 1998 LeNet-5 (early CNN) achieved 99.2% on MNIST - a breakthrough showing spatial processing matters!